



Vulnerability Detection - Fuzzing

Holistic Software Security

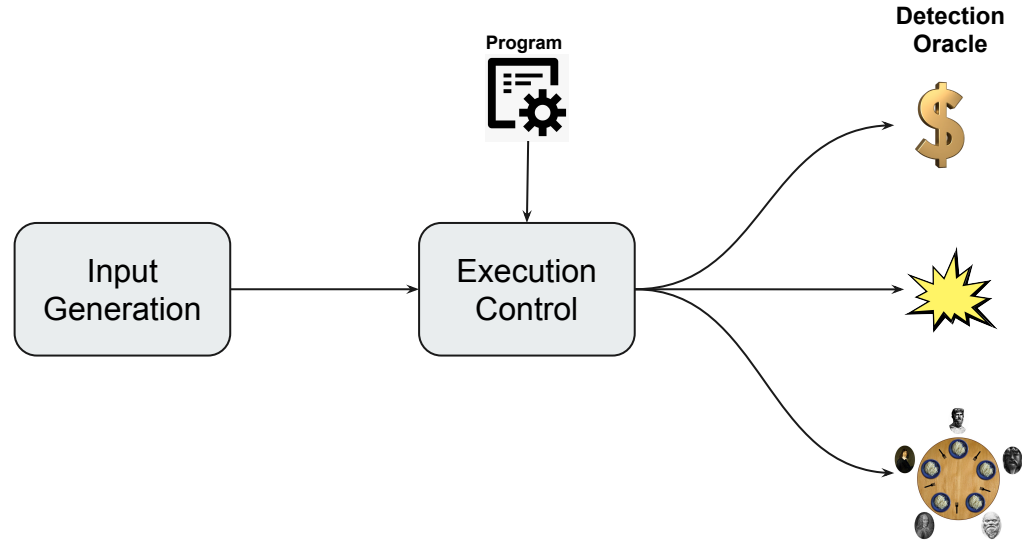
Aravind Machiry



Fuzzing

- Automated test generation using random data.
 - Generate effective test cases, primarily using random data.

Fuzzing: High Level Idea





Input Generation

- Generate inputs (mostly randomized) to be fed into the program:
 - Random source.
 - Mutating existing inputs.
 - Based on a given input grammar.



Execution Control

- Execute the program with a given input:
 - Regular command line programs: `execve` and `stdin`.
 - OS: System calls.
 - Network programs: Send over network.
 - Input file: Save the data into a file and provide file.



Detection Oracle

- Detection of interesting program behavior:
 - Program crash.
 - Race condition.
 - High execution time.

Fuzzing Success

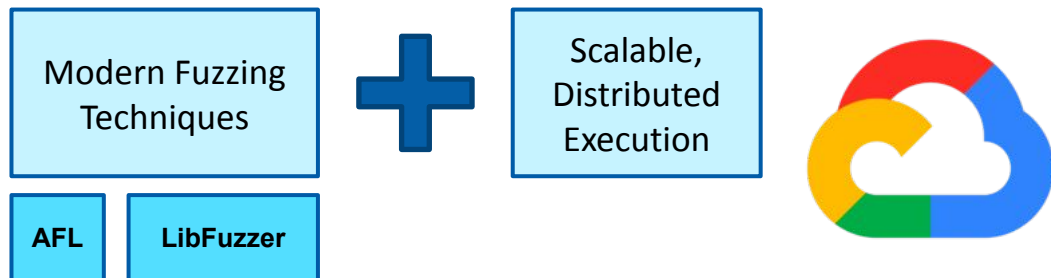
GCC Bug List Found by Random Testing (Total 79)									
date open	bug_id*	bug type	priority	rev reported	platform	component	status	date fixed(rev)(by)	File modified (lines)
3/30/2008	35764	wrong	P3	4.3.0	x86-32	target	confirmed	n/a	
5/15/2008	36238	crash	P2	4.4.0	x86-32	target	fixed	08/10 138924(Pinski)	reload1.c(1)
6/17/2008	36548	wrong	P3	136854	x86-32	middle-end	fixed	08/22 139450(Guenther)	fold-const.c(12)
6/24/2008	36613	wrong	P1	137045	x86-32	target	fixed	08/11 138955(Matze)	reload1.c(8)
6/25/2008	36635	crash	P1	137122	x86-32	target	fixed	10/08 140966(Jelinek)	cse.c(11)
7/1/2008	36691	wrong	P1	137327	x86-32	middle-end	fixed	08/04 138645(Guenther)	tree-ssa-loop-niter.c(2)
8/13/2008	37102	wrong	P1	139046	x86-32	tree-opt	fixed	10/17 141195(Macleod)	tree-outof-ssa.c(95)
8/13/2008	37103	wrong	P3	139048	x86-32	middle-end	fixed	08/14 139004(Jelinek)	fold-const.c(11)

open (83/7)							
Title	Repro	Cause bisect	Fix bisect	Count	Last	Reported	Last activity
BUG: scheduling while atomic: syz-executor/ADDR	C	done		1	4401h	1b08a	1b08a
BUG: sleeping function called from invalid context in __fput				1	4401h	1b39a	1b38a
UBSAN: shift-out-of-bounds in init_sh				1	2d04h	3b19a	3b19a
BUG: sleeping function called from invalid context in __fdget_pos				1	6d06h	2d00b	7b41a
unexpected kernel reboot (6)				1	2d03h	2d02h	2d02h
INFO: task can't die in p9_client_rpc(3)				4	1d01h	2d13h	2d13h
memory leak in j1939_sk_sendmsg	C			1	6d15h	2d15h	2d05h
KASAN: use-after-free Read in v4l2_ioctl(2)	C	error		1	9d14h	5d14h	4d11h
KASAN: out-of-bounds Read in do_exit				1	10d	6d14h	4d14h
memory leak in xfrm_user_rcv_msg	C			1	12d	8d00h	12b54a
BUG: corrupted list in kobject_add_internal(3)		inconclusive		1	12d	8d01h	6d06h
memory leak in j1939_xtp_rx_rts	syz			1	12d	8d02h	5d09h
INFO: task hung in port100_probe	C	error		3	12d	8d04h	8d03h
general protection fault in detach_extent_buffer_page				1	14d	9d14h	9d10h



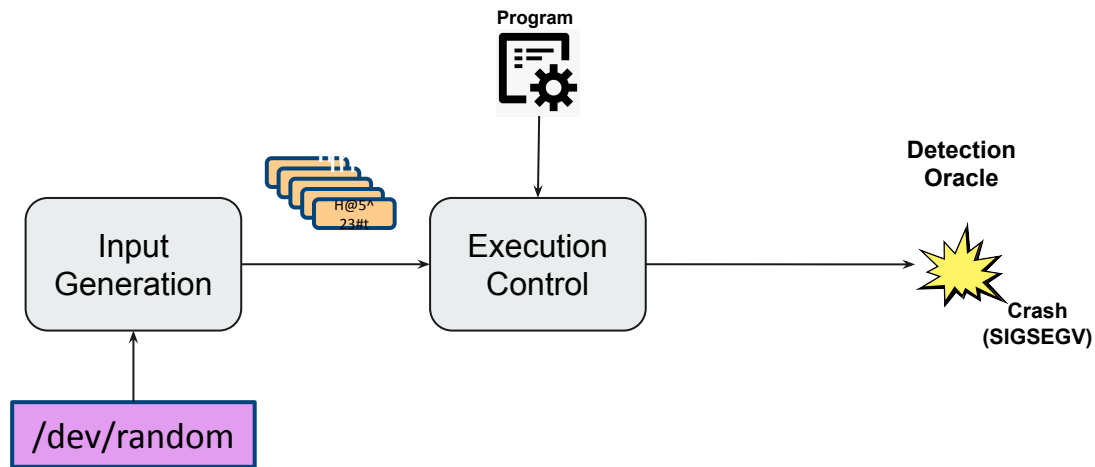
Fuzzing Success: OSS-Fuzz

- Continuous fuzzing infrastructure hosted on the [Google Cloud Platform](#)



- OSS-Fuzz has discovered over 17,400 bugs from 2016 to 2019 in many large projects (e.g. openssl, llvm, postgresql, git, firefox)

Fuzzing: Gen 1 (Random data)





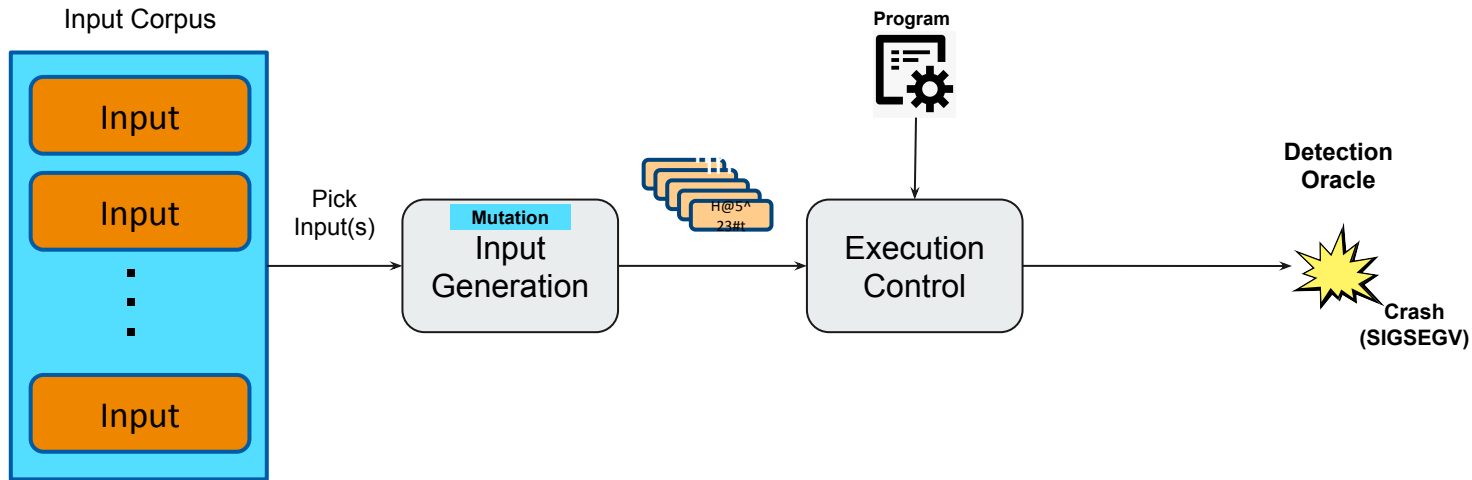
Fuzzing: Gen 1

- Conducted by Barton Miller @ Univ of Wisconsin.
- 1990: Command-line fuzzer, testing reliability of UNIX programs.
 - Bombards utilities with random data
- 1995: Expanded to GUI-based programs (X Windows), network protocols, and system library APIs.
- Later: Command-line and GUI-based Windows and OS X apps.

Caused 25-33% of UNIX utility programs to crash (dump state) or hang (loop indefinitely).

- Hard to generate well formed data:
 - E.g., PNG files.

Fuzzing: Gen 2.a (Mutation based)

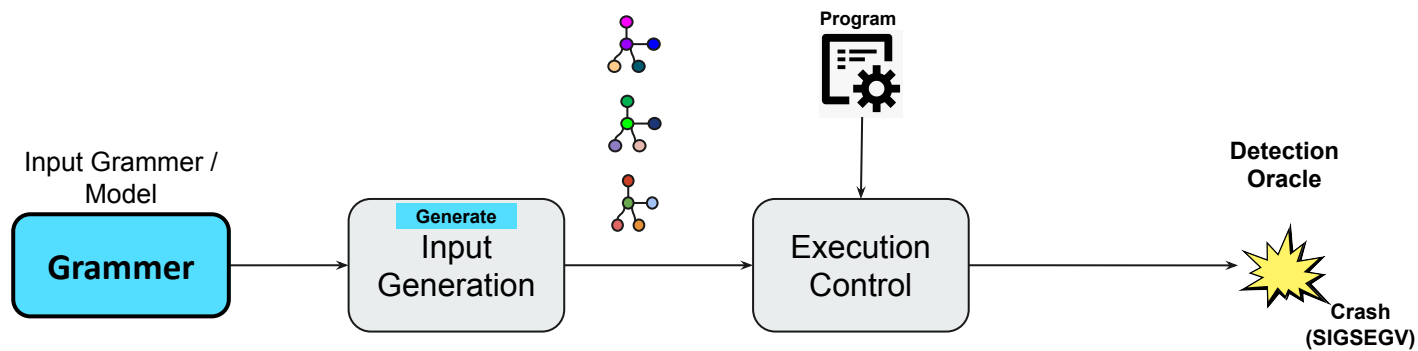




Fuzzing: Gen 2.a


- Very effective at generating semi-structured inputs.
- Still not so effective at generating highly structured inputs:
 - E.g., C files.

Fuzzing: Gen 2.b (Generation based)

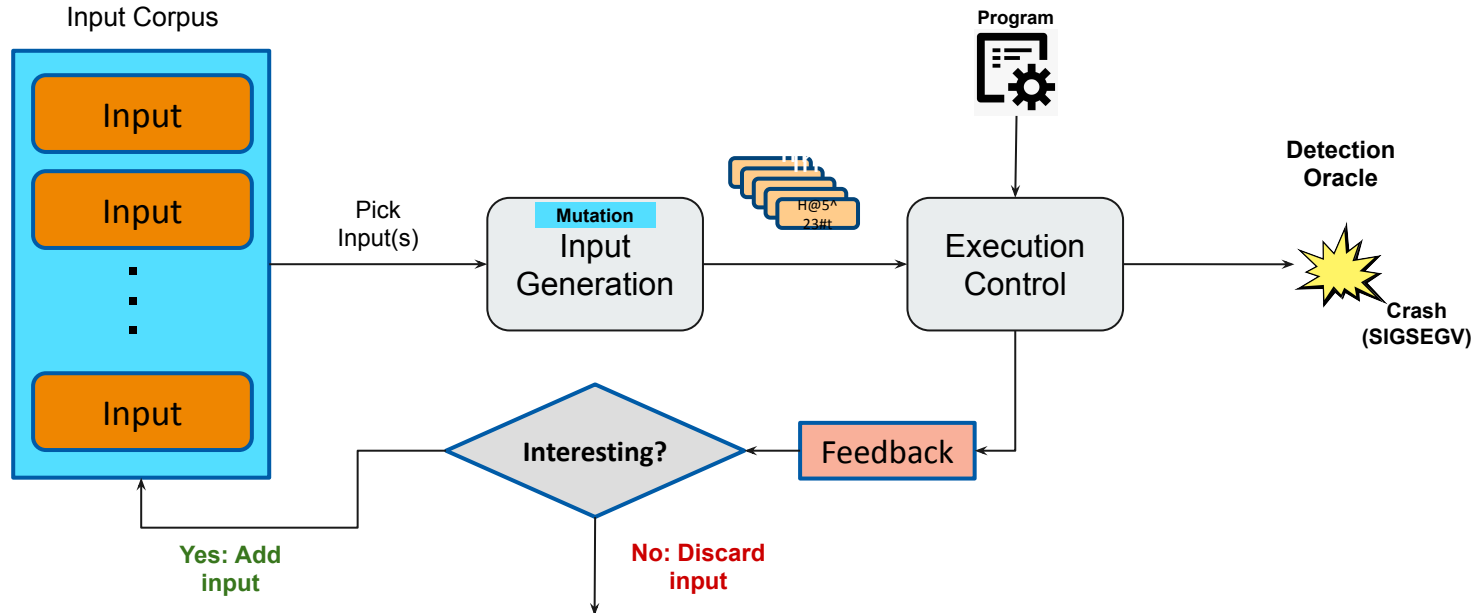




Fuzzing: Gen 2.b

- Very effective at generating complex inputs:
 - Csmith: Generate syntactically valid but random C programs.
- Commercial tools:
 - Peach. 
- Need to manually write these input grammars:
 - Domain Specific Language.
 - Large: ~200 lines

Fuzzing: Gen 3 (Feedback guided Mutation based)





Fuzzing: Gen 3

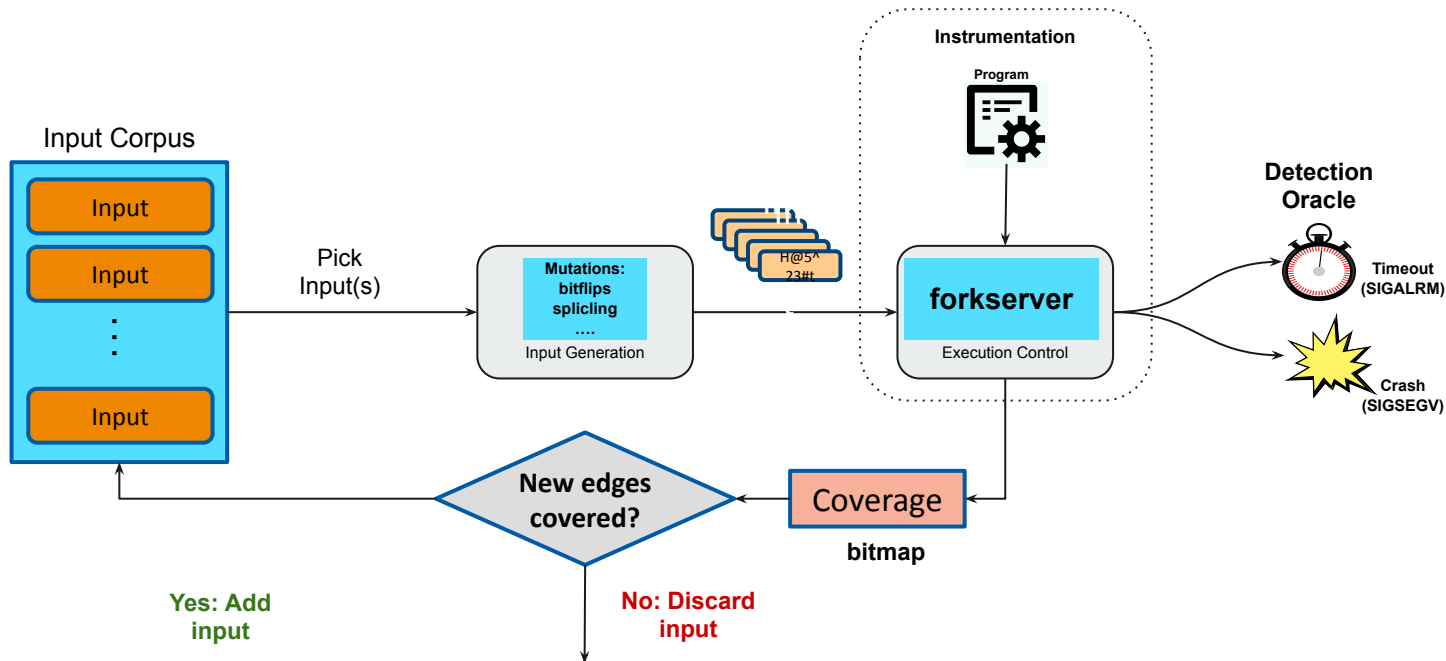
- Extremely effective at quickly generating well-formed inputs.
- Highly successful commercial grade tool:
 - AFL (AFLPlusPlus)
- Need a way to capture feedback: Impacts performance.



Fuzzer Deep Dive : AFLPlusPlus

- Based on American Fuzzy Lop (AFL) developed by Michał Zalewski
- Coverage Feedback based Mutational Fuzzing.
- Highly customizable, efficient and very well maintained.
- Revolutionized fuzzing research:
 - ~30 papers since 2015.
- Found various (~200) bugs in well-maintained programs.

AFLPlusPlus (A++): Coverage guided





A++: Coverage Map

- Coverage choices:
 - Line or Basic block coverage:
 - Basic blocks executed.
 - Edge coverage (used by A++):
 - Edges (Basic block tuple) executed.
 - Path coverage:
 - Sequence of basic blocks executed.

Precision




Overhead




Coverage choices

Execution 1



Execution 2



- Does Execution 1 and 2 have:
 - Same basic block coverage?
 - Same edge coverage?
 - Same path coverage?




A++: Coverage Instrumentation



- Coverage map: Memory area in the program that stores coverage.
- Every basic block in all functions will be instrumented to update coverage map.

```
cur_location = <COMPILE_TIME_RANDOM>;  
coverage_map[cur_location ^ prev_location]++;  
prev_location = cur_location >> 1;
```



(Hopefully) Unique key
for each edge




A++: Coverage Map




A++: Coverage Map




A++: Coverage Map: Bucketized edge counts




- Edge counts are bucketized:
 - E.g., Coverage map of executions with loop counts that belong to the same bucket will be considered the same.



A++: Coverage Map



**Has same
coverage
map as**



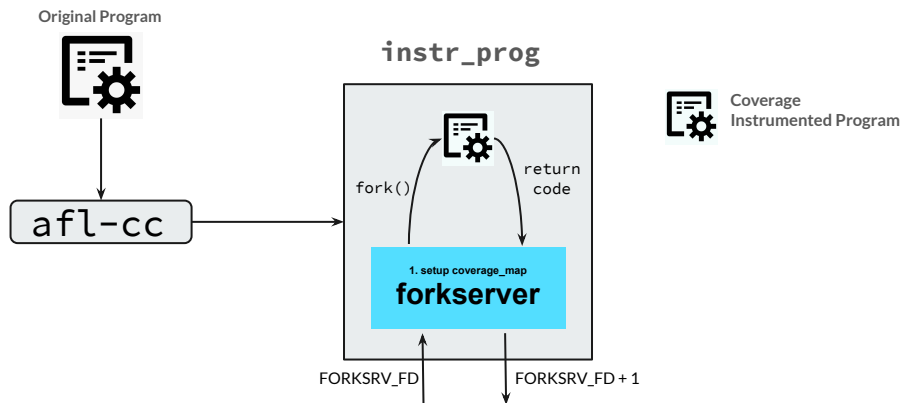
Using A++

- Instrumentation: Compile the target program using afl compiler, i.e., afl-cc:
 - afl-cc <.c> -o instr_prog
 - Does:
 - Instrumentation to compute coverage.
 - Add forkserver.
- Fuzzing: Start fuzzing instr_prog:
 - afl-fuzz -i <inputs_folder> -o <output_folder> -- instr_prog

```
american fuzzy lop ++3.13a (default) [fast] {0}
- process timing -
  run time : 0 days, 0 hrs, 0 min, 42 sec
  last new path : 0 days, 0 hrs, 0 min, 2 sec
  last uniq crash : none seen yet
  last uniq hang : none seen yet
- cycle progress -
  now processing : 13.6 (76.5%)
  paths timed out : 0 (0.00%)
- stage progress -
  now trying : havoc
  stage execs : 6990/9418 (74.22%)
  total execs : 417k
  exec speed : 10.1k/sec
- fuzzing strategy yields -
  bit flips : disabled (default, enable with -D)
  byte flips : disabled (default, enable with -D)
  arithmetics : disabled (default, enable with -D)
  known ints : disabled (default, enable with -D)
  dictionary : n/a
  havoc/splice : 16/394k, 0/15.1k
  py/custom/rq : unused, unused, unused
  trim/eff : 0.00%/5, disabled
- overall results -
  cycles done : 260
  total paths : 17
  uniq crashes : 0
  uniq hangs : 0
- map coverage -
  map density : 3.12% / 7.81%
  count coverage : 8.00 bits/tuple
- findings in depth -
  favored paths : 2 (11.76%)
  new edges on : 17 (100.00%)
  total crashes : 0 (0 unique)
  total tmouts : 0 (0 unique)
- path geometry -
  levels : 7
  pending : 9
  pend fav : 0
  own finds : 16
  imported : 0
  stability : 0.00%
[cpu000: 33%]
```

A++: Instrumentation

```
afl-cc <.c> -o instr_prog
```



- `forkserver (constructor) afl-compiler-rt.o.c:`

Setup coverage map (shared memory map).

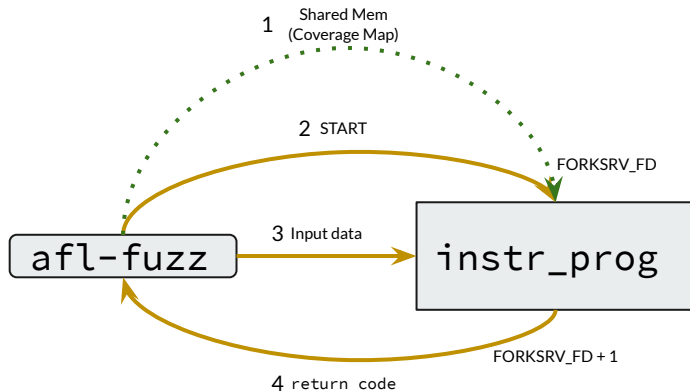
```
while (1){
```

1. Wait for command at `FORKSRV_FD`.
2. Once received, fork and start executing main of the original program:
 - Program would be writing to coverage map (shared memory).
3. Sends the return code through `FORKSRV_FD + 1`

```
}
```

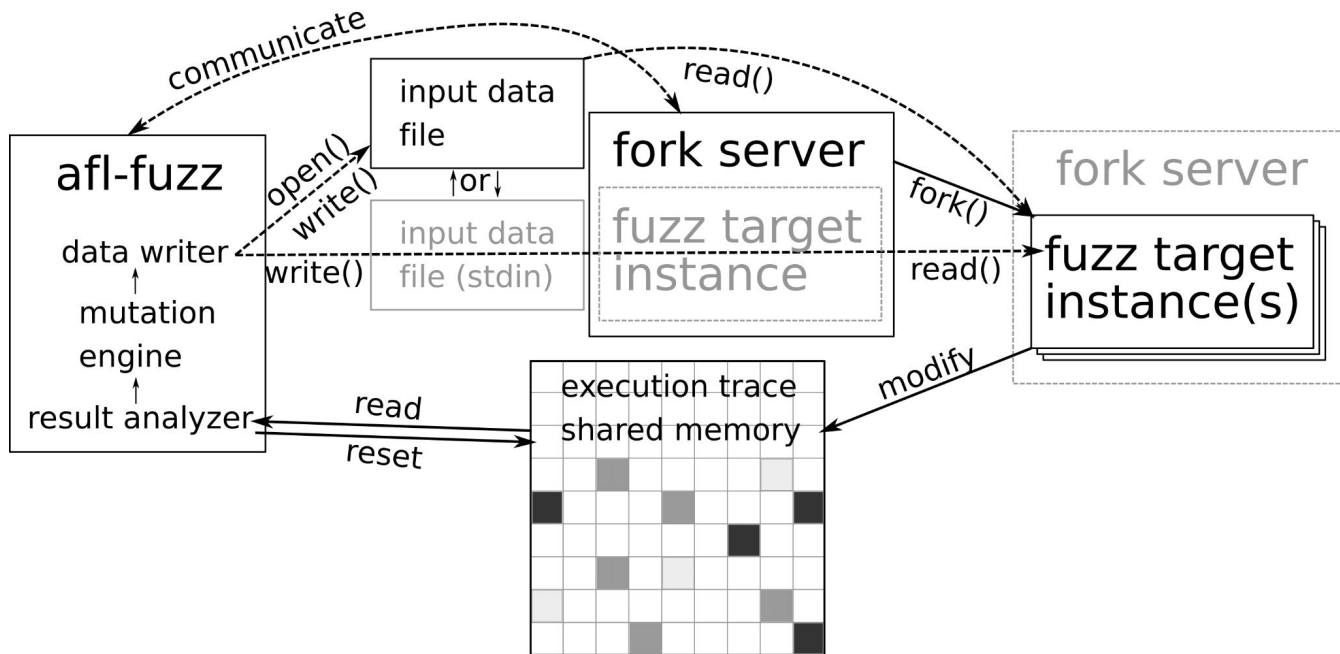
A++: Fuzzing

```
afl-fuzz -i <inputs_folder> -o <output_folder> -- instr_prog
```



- `afl-fuzz (src/afl-fuzz.c):`
 - Setup (1):
 - Send shared memory id.
 - Fuzzing Loop (2-3-4):
 - 2. START
 - 3. Input data (stdin or file)
 - 4. return code (crashes or timeout)

A++: Fuzzing





A++: Few drawbacks

- Effectiveness highly depends on the quality of initial test cases.
- Does not readily accepts grammar for inputs.
- Does not readily accepts other coverage metrics:
 - We may want different coverage metric for functions:
 - E.g., BB coverage for foo, edge for bar, path for baz.



Fuzzing Challenges: Input Generation

- Constrained Input:

- Driller: Augmenting Fuzzing through Symbolic Execution [NDSS 2016]
- Angora: Efficient Fuzzing by Principled Search [S&P 2018]
- REDQUEEN: Fuzzing with Input-to-State Correspondence [NDSS 2019]

```
if (i == 345890)  
{  
    ...  
}
```

- Structured Input:

- DIFUZE: Interface Aware Fuzzing for Kernel Drivers [CCS 2017]
- WEIZZ: Automatic Grey-Box Fuzzing for Structured Binary Formats [ISSTA 20]

```
1 typedef struct {  
2     ISP_RT_BUF_CTRL_ENUM ctrl;  
3     _isp_dma_enum_ buf_id;  
4     ISP_RT_BUF_INFO_STRUCT *data_ptr;  
5     ISP_RT_BUF_INFO_STRUCT *ex_data_ptr;  
6     unsigned char *pExtend;  
7 } ISP_BUFFER_CTRL_STRUCT;
```



Fuzzing Challenges: Coverage metrics

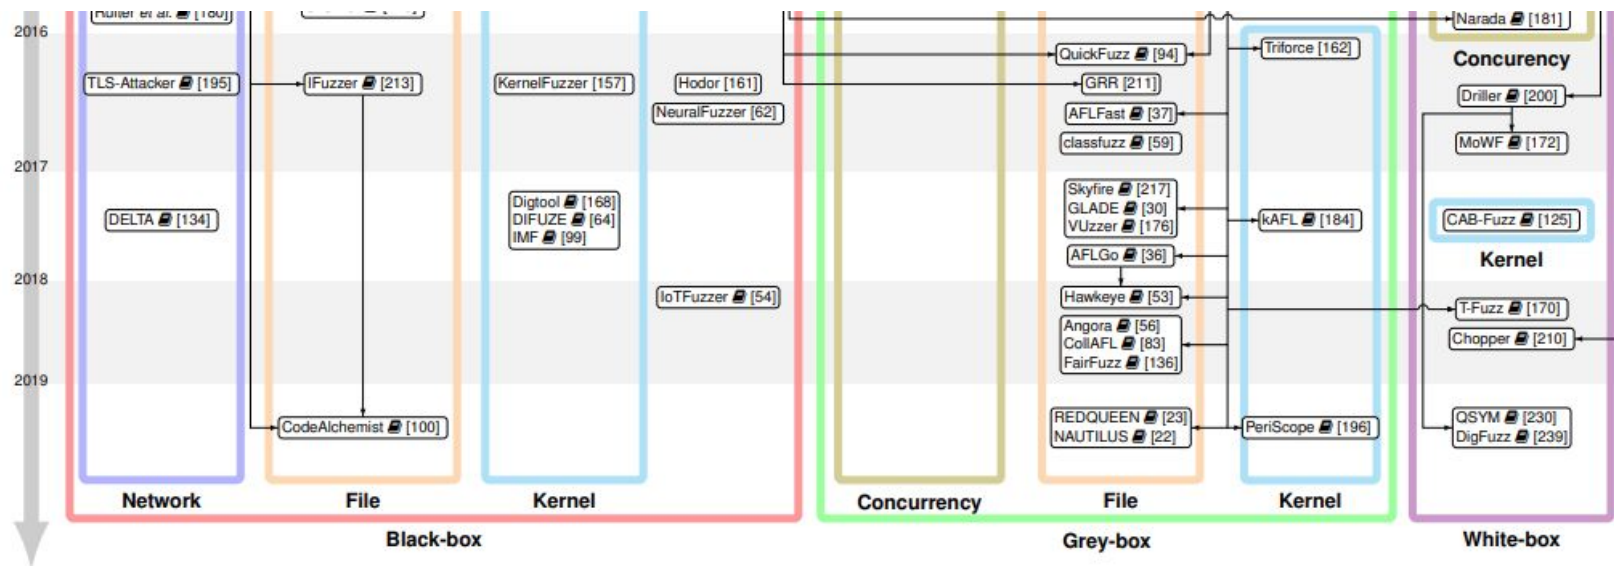
- Is Path Coverage always good?
 - “Be Sensitive and Collaborative: Analyzing Impact of Coverage Metrics in Greybox Fuzzing” [RAID 2019]
 - “CollAFL: Path-Sensitive Fuzzing” [S&P 18]



Fuzzing Challenges: Input prioritization

- Some inputs are good than other inputs?
 - “Not All Coverage Measurements Are Equal: Fuzzing by Coverage Accounting for Input Prioritization” [NDSS 2020]
 - “ParmeSan: Sanitizer-guided Greybox Fuzzing” [USENIX 2020]

Fuzzing Trends





Fuzzing Trends

- New directions:
 - ML to detect which bytes to mutate.
 - Transform the program and make it easy to fuzz (t-fuzz).
 - Combine different fuzzers: CollabFuzz: A Framework for Collaborative Fuzzing [EuroSec 2021]
- Improvements:
 - Use fancy techniques to improve different aspects of fuzzing.
- Fuzzing different applications:
 - File systems, Kernel drivers, IoT devices, etc.



Fuzzing as a generic exploratory technique

- Fuzzing allows us to find inputs that has high probability to satisfy certain goal.
 - Goal: Find more bugs:
 - Feedback: Coverage.
 - Goal: Find more temporal bugs (e.g., use-after-free, double-free):
 - Feedback: Likelihood of input triggering malloc/free.
 - Goal: Find concurrency bugs.
 - Feedback: Number of threads invoked.
 - Goal: Find denial-of-service bugs.
 - Feedback: Time taken by the input (the more time the better).
 - Goal: Type inference: Infer types for variables.
 - Feedback: Number of type-checker error (the less the better).