



# Vulnerability Detection - Sanitizers

Holistic Software Security

Aravind Machiry



# Bug Manifestations

- How bugs affect program behavior?
  - If we have exhaustive test cases:
    - Actual output != Expected output.
  - In the absence of test cases, i.e., Fuzzing:
    - Memory errors: Program Crashes (SIGSEGV) => Access/Execute invalid memory.
    - There could be bugs which do not result in SIGSEGV.



# Silent Bugs

```
int main() {  
    unsigned i, j, a[2];  
    scanf("%u %u", &i, &j);  
    a[i] = j;  
    ...  
    return 0;  
}
```

Here, if  $i == 2$  (off by one error), the program may not crash?

- Why?

## Runtime Stack

Return Address
i
j
unsigned a[2]

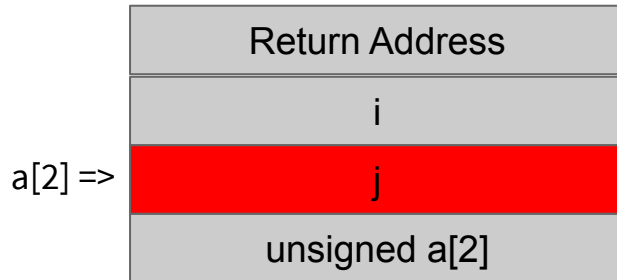
# Silent Bugs

```
int main() {  
    unsigned i, j, a[2];  
    scanf("%u %u", &i, &j);  
    a[i] = j;  
    ...  
    return 0;  
}
```

Here, if  $i == 2$  (off by one error), the program may not crash?

- Why?

## Runtime Stack





# Improving bug detection

- The behavior of a bug, especially memory corruption, depends on the program state and execution environment.
- Can we detect these bugs without relying on program state?
  - Fuzzing: we detect a bug if it results in the program crash (SIGSEGV).
  - Idea: Make all bugs result in program crashes.



# Sanitizers

- Change the program such that *we detect bugs when they occur* instead of waiting for the bugs to result in crash.
- Mechanism: Instrument the program by adding additional checks for detecting bugs.

# Sanitizers: Overview

Original Program



Instrumentation

Instrumented Program



Program with  
additional Checks

Original Program

```
int main() {  
    unsigned i, j, a[2];  
    scanf("%u %u", &i, &j);  
    a[i] = j;  
    ...  
    return 0;  
}
```



Array out-of-bounds Sanitizer

```
int main() {  
    unsigned i, j, a[2];  
    scanf("%u %u", &i, &j);  
    if (i < 2) {  
        a[i] = j;  
    } else { CRASH}  
    ...  
    return 0;  
}
```

Instrumented Program

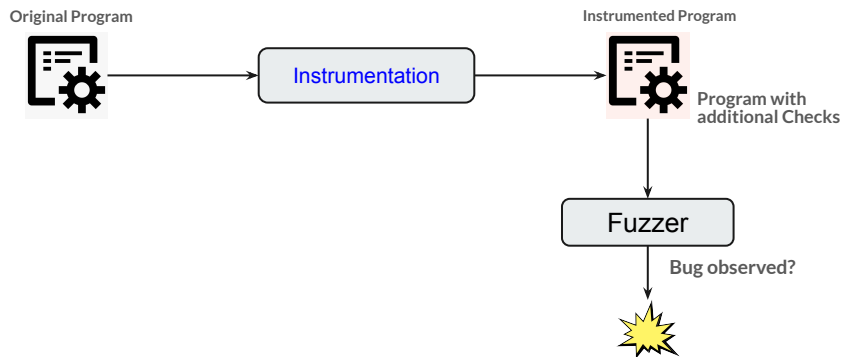


# Real world Sanitizers

- Usually bug specific. Examples:
  - MemorySanitizer: Detects *uninitialized reads*.
  - AddressSanitizer: Detects *invalid memory accesses*.
- **General instrumentation idea:** At all instructions in the program where the bug can occur, add a check to detect the bug.
  - AddressSanitizer: Detects Invalid Memory Accesses.
    - Invalid Memory access can occur at load and store instructions.
      - **Instrument every load and store to check if the used address is invalid (i.e., does not belong to a program object).**



# Sanitizers: Usage



Original Program

```
int main() {  
    unsigned i, j, a[2];  
    scanf("%u %u", &i, &j);  
    a[i] = j;  
    ...  
    return 0;  
}
```

Array out-of-bounds Sanitizer

```
int main() {  
    unsigned i, j, a[2];  
    scanf("%u %u", &i, &j);  
    if (i < 2) {  
        a[i] = j;  
    } else { CRASH}  
    ...  
    return 0;  
}
```

Instrumented Program



# Why can't we always use sanitizers?

They detect bugs at runtime => Why can't we just use sanitizers and not worry about bugs, as they will never lead to vulnerabilities.

**THERE'S  
NO  
FREE  
LUNCH**



Sanitizers introduce **a lot** of overhead.



# Sanitizers Implementation

- Sanitizers need to maintain lot of additional state to check for the possibility of bugs.
  - AddressSanitizer: Detects Invalid Memory Accesses:
    - Need to **maintain metadata** regarding which memory (i.e., address) is valid v/s invalid.
    - Tricky: Handling dynamic memory allocation.

**Popular research direction:** Smart and efficient way to maintain metadata.

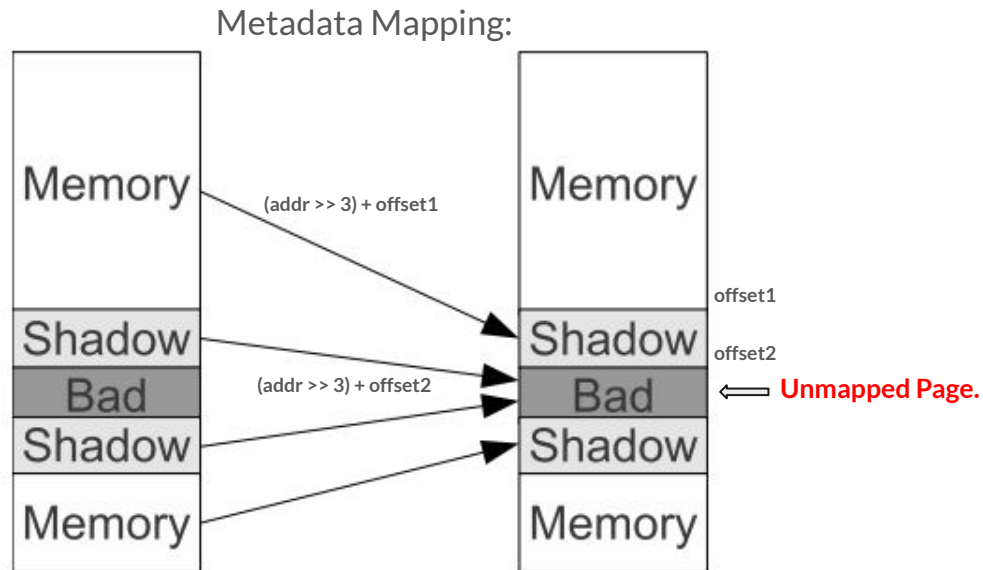


# AddressSanitizer (ASan)

- Metadata (or shadow memory):
  - One eighth of the virtual memory will be used to maintain metadata:
    - One bit of metadata for each byte of application memory.
    - Bit is zero: The corresponding address is valid else invalid.
- Accessing metadata for a given address (Addr):
  - Direct Mapping:

```
// Checking 8-byte access
MetadataAddr = (Addr >> 3) + Offset;
if (*MetadataAddr != 0)
    ReportAndCrash(Addr);
```

# ASan: Mapping





# ASan: Usage

a.c

```
void foo(T *a) {  
    *a = 0x1234;  
}
```



clang -fsanitize=address a.c -c -DT=long



```
push %rax  
mov %rdi,%rax  
shr $0x3,%rax  
mov $0x1000000000000,%rcx  
or %rax,%rcx  
cmpb $0x0,(%rcx) # Compare Shadow with 0  
jne 23 <foo+0x23> # To Error  
movq $0x1234,(%rdi) # Original store  
pop %rax  
retq  
callq __asan_report_store8 # Error
```



# ASan: Conclusion

- One of the most popular sanitizers: Used extensively in fuzzing.
- Overhead:
  - Adds additional instructions:
    - Memory overhead: **~3X** (Consumes thrice the amount of memory).
    - Slowdown: **~2X** (Runs at half the speed).





# ThreadSanitizer

- Detects data races.
- Where can data races happen i.e., which instructions it should track?
- How to detect a data race? What metadata should be maintained?



## Other sanitizers (supported by clang)

- `-fsanitize=address`: AddressSanitizer, a memory error detector.
- `-fsanitize=thread`: ThreadSanitizer, a data race detector.
- `-fsanitize=memory`: MemorySanitizer, a detector of uninitialized reads. Requires instrumentation of all program code.
- `-fsanitize=undefined`: UndefinedBehaviorSanitizer, a fast and compatible undefined behavior checker.
- `-fsanitize=dataflow`: DataFlowSanitizer, a general data flow analysis.



# Sanitizers: Final Thoughts

- They increase the ability of fuzzing to find bugs.
- Always use them with fuzzers: Performance impact does not matter much - lets throw more machines.
- New sanitizers => Always appreciated and could have a high impact.
- Decreasing overhead of sanitizers: Appreciated but may have less impact.